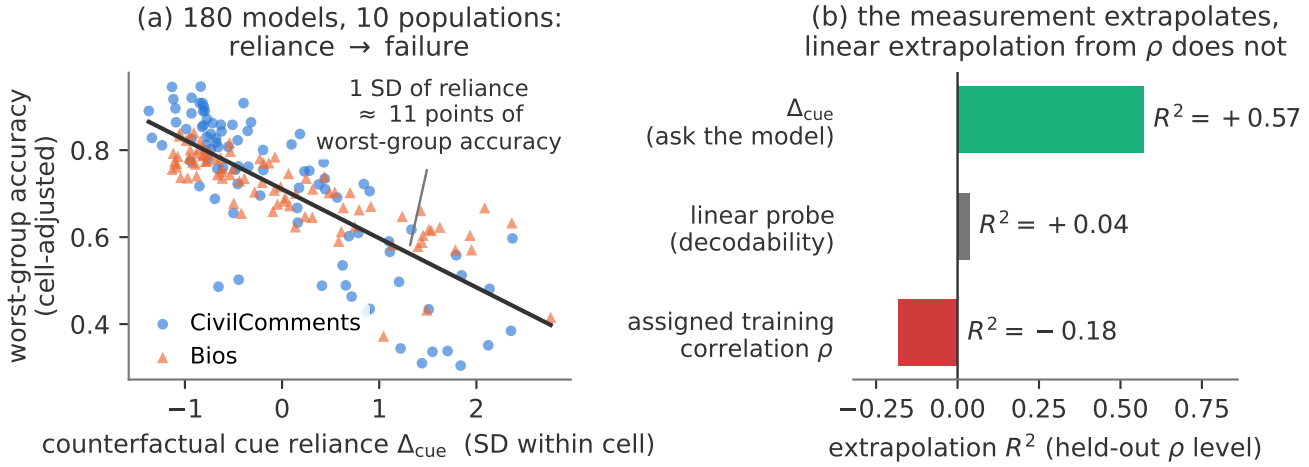
(a) 180 models, 10 populations:
reliance → failure
1 SD of reliance
≈ 11 points of
worst-group accuracy
worst-group accuracy
(cell-adjusted)
0.8
0.6
0.4
CivilComments
Bios
−1
0
1
2
counterfactual cue reliance Δ_{cue} (SD within cell)
(b) the measurement extrapolates,
linear extrapolation from ρ does not
Δ_{cue}
(ask the model)
$R^2 = +0.57$
linear probe
(decodability)
$R^2 = +0.04$
assigned training
correlation ρ
$R^2 = -0.18$
−0.25
0.00
0.25
0.50
0.75
extrapolation R^2 (held-out ρ level)